



San Francisco

VULKAN

A META DATA DRIVEN
BACKUP RESTORE SYSTEM

INTRODUCTION - VULKAN

VULKAN - DISTRIBUTION OF SUBSETS OF DATA, BETWEEN DATABASE SYSTEMS, IN A LIGHTWEIGHT, CONFIGURABLE AND SCALABLE FASHION.

PRESENTATION CONTEXT

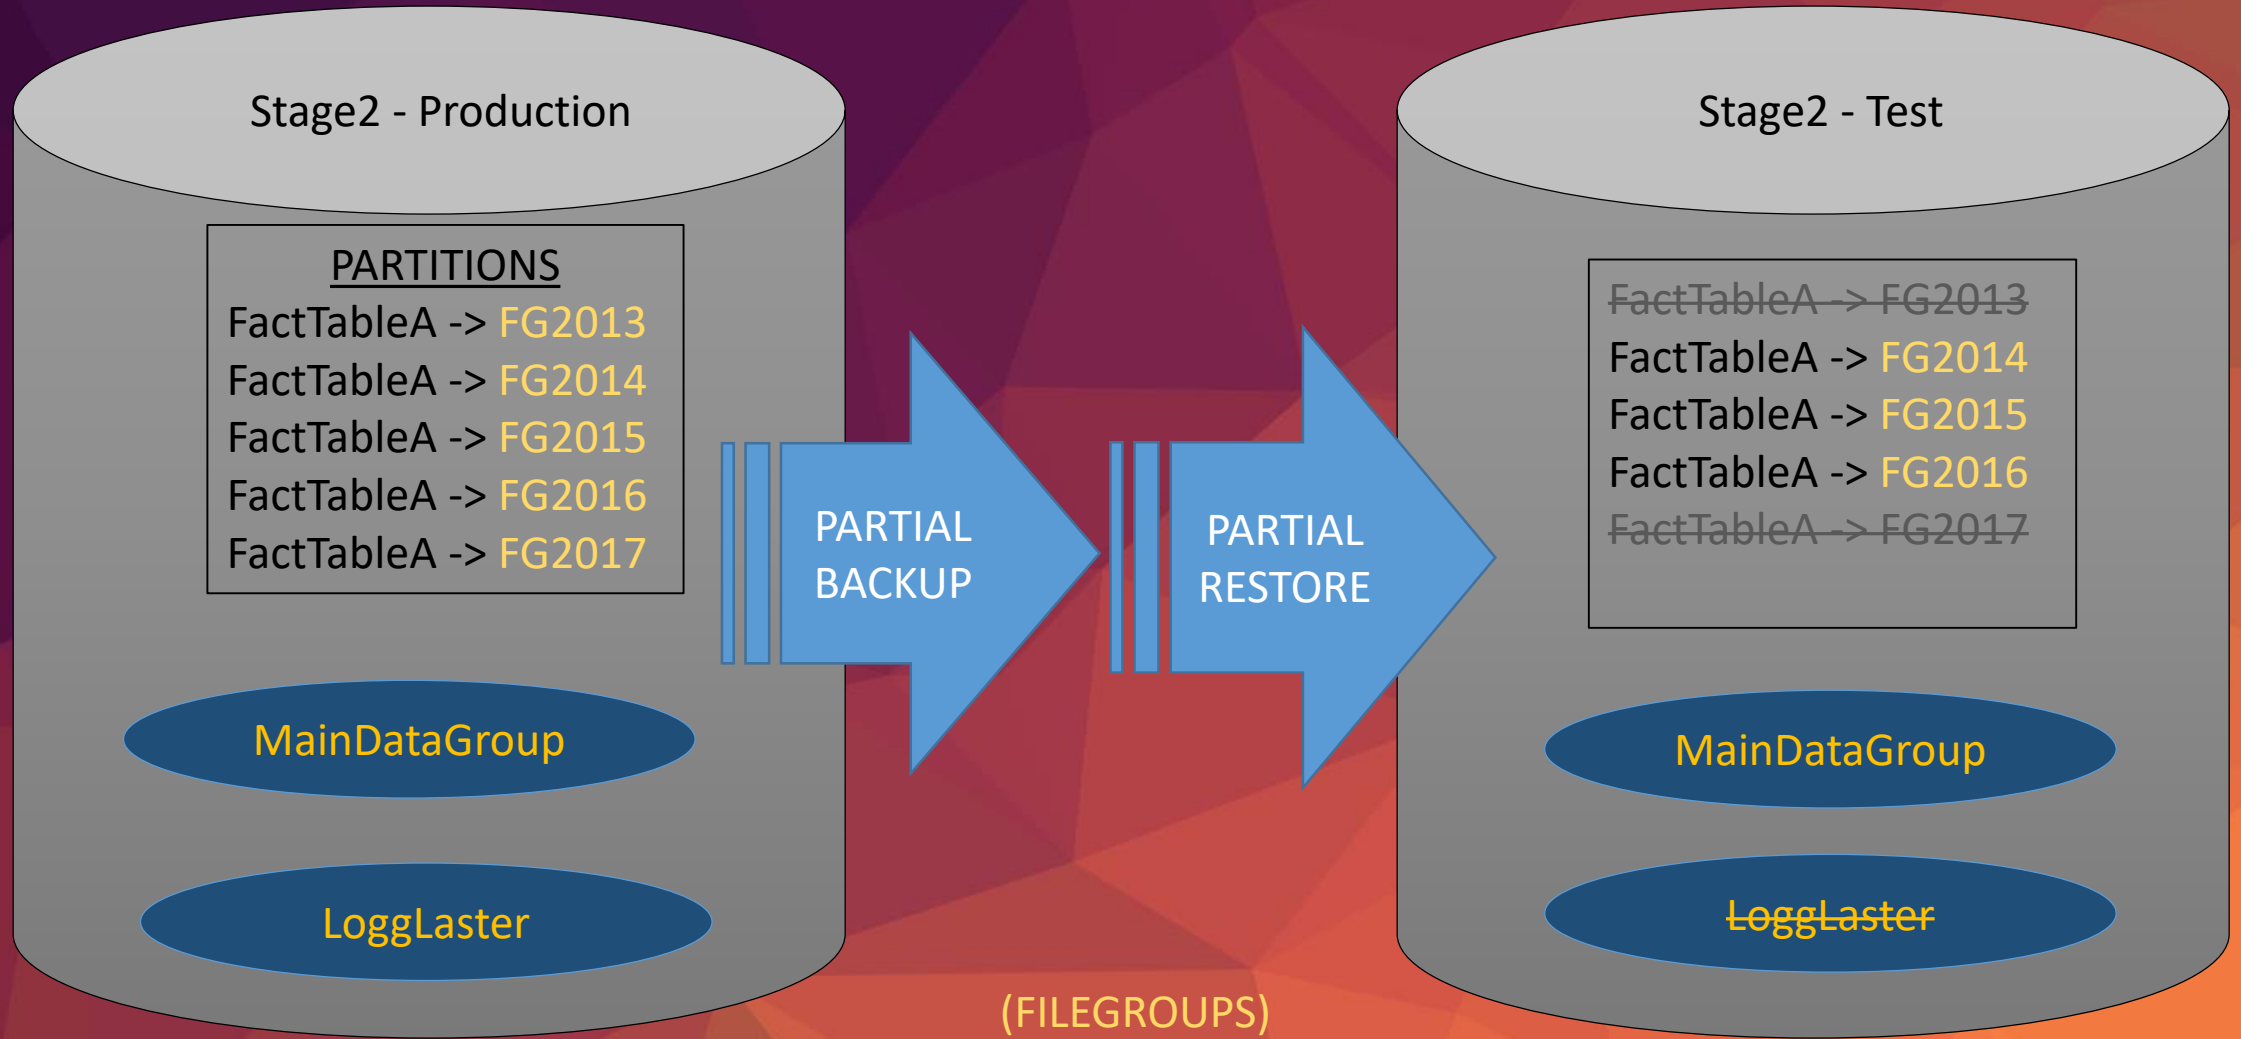
- DATAWAREHOUSING
- SQL SERVER
- BACKUP / RESTORE OF DATABASES
- FILEGROUPS = CONTAINERS OF DATA (TABLES / INDEXES)
- PARTITIONING OF DATA
- META DATA DRIVEN SYSTEMS
- BIML (Business Intelligence Markup Language)
- INTEGRATION SERVICES (SSIS)

BACKGROUP – MISSION STATEMENT

“I want us to be able to distribute subsets of production data, to other subsystems, in a light-weight, configurable and scalable fashion”

In the context of distrubiting subsets of data, we want:

- Ability to completely include/omit entire islands of data. *(For example, omit large logg tables)*
 - Vertical Scalability
- Ability to include/omit specific levels of data, i.e. Include/omit ranges of data, usually time ranges, within **individual** tables. *(For example include only last 3 years of data)*
 - Horizontal Scalability



Vulkan – A Technical Overview

- ❖ Handles Full or Partial Distribution of subsets of data, to SubSystems (Typical Production Data to Test Systems)
- ❖ System Parameters in Configurable Meta Data SQL Tables
- ❖ SSIS packages to execute all parts of system workflow
- ❖ BIML Code, in conjunction with Metadata, to Auto generate SSIS Packages

TERMINOLOGY - DISTRIBUTION OF DATA USING SQL SERVER

BACKUP – Backup a full or partial database to a *.BAK-file

RESTORE – Restore a database from a *.BAK-file, Full or Partial

FILEGROUPS – A way to group set of data, logical (and physical)

DATAFILE – The physical file of data, attached to a certain FILEGROUP

CLUSTERED INDEX – Instructions (Function) on how and where a table is stored on disk

PARTITIONS – A way to split tables into subsections, logical (and physical), based on fixed ranges

MERGE - Merging of Partitions

SWITCH – Switching of data, in and out of Partitions, using Shadow-tables

PREPARING FOR DATA DISTRIBUTION – SET UP PARTITIONING

Sample table: dbo.FactTablePeriode (Partition Column: PERIOD (DataType INT))

Step 1. Create FileGroups and DataFiles, as containers for Partitioned Data (FG2014 -> FG2014.ndf...)

Step 2. Create Partition Function to support Range and DataType

```
CREATE PARTITION FUNCTION [IntRange_pf](INT) AS RANGE RIGHT FOR VALUES
```

```
(  
20030101, 20040101, 20050101, 20060101, 20070101, 20080101,  
20090101, 20100101, 20110101, 20120101, 20130101, 20140101,  
20150101, 20160101, 20170101, 20180101, 20190101, 20200101,  
20210101, 20220101, 20230101, 20240101, 20250101  
)
```

RANGE RIGHT; example: A row with PERIODE value 20060101, will be placed to the partition to the RIGHT (FG2006)

RANGE LEFT; example: A row with PERIODE value 20060101, will be placed to the partition to the LEFT (FG2006)

PREPARING FOR DATA DISTRIBUTION – SET UP PARTITIONING

Sample table: dbo.FactTablePeriode (Partition Column: PERIOD (DataType INT))

Step 3. Create Partition Schema, to align Ranges with correct FileGroup (and Data File)

```
CREATE PARTITION SCHEME IntScheme_ps  
AS PARTITION [IntRange_pf]  
TO (
```

```
    FG2002,FG2003,FG2004,FG2005,FG2006,  
    FG2007,FG2008,FG2009,FG2010,FG2011,  
    FG2012,FG2013,FG2014,FG2015,FG2016,  
    FG2017,FG2018,FG2019,FG2020,FG2021,  
    FG2022,FG2023,FG2024,FG2025
```

```
)
```

PREPARING FOR DATA DISTRIBUTION – SET UP PARTITIONING

Sample table: dbo.FactTablePeriod (Partition Column: PERIOD (DataType INT))

Step 4. Move table data into Partitons with Aligned Clustred Index, with DROP_EXISTING on Partition Scheme

```
CREATE UNIQUE CLUSTERED INDEX [ClusteredIndexPeriod] ON [dbo].[FactTablePeriod]
(
    [Period] ASC,
    [SalesOrderID] ASC,
    [SalesOrderDetailID] ASC
)WITH (PAD_INDEX = OFF, STATISTICS_NORECOMPUTE = OFF, SORT_IN_TEMPDB = OFF, IGNORE_DUP_KEY = OFF,
DROP_EXISTING = ON, ONLINE = OFF, ALLOW_ROW_LOCKS = ON, ALLOW_PAGE_LOCKS = ON)
ON [IntScheme_ps]([Period])
```



BACKUP DATABASE PARTIAL

Sample table: dbo.FactTablePeriode (Partition Column: PERIOD (DataType INT))

Step 5. Backup Database WITH PARTIAL (in Full Recovery Mode)

BACKUP DATABASE [Stage2]

```
FILEGROUP = 'PRIMARY',  
FILEGROUP = 'MainDataGroup',  
FILEGROUP = 'FG2014',  
FILEGROUP = 'MainIndexGroup',  
FILEGROUP = 'EVRYGroup',  
FILEGROUP = 'UDBGroup',  
FILEGROUP = 'FG2015',  
FILEGROUP = 'FG2016'
```

```
TO DISK = 'E:\SQLBackup\BackupFileforRestoreToTest\Stage2.bak'
```

```
WITH COMPRESSION,  
COPY_ONLY,  
NOFORMAT,  
NOINIT,  
NAME = N'Stage2 Backup',  
SKIP,  
REWIND,  
NOUNLOAD;
```

RESTORE DATABASE PARTIAL

Sample table: dbo.FactTablePeriode (Partition Column: PERIOD (DataType INT))

Step 6. Restore Database on Destination Server WITH PARTIAL (in Full Recovery Mode)

```
RESTORE DATABASE [VulkanProd]
```

```
FILEGROUP = 'FG2015',  
FILEGROUP = 'FG2016',  
FILEGROUP = 'FG2017',  
FILEGROUP = 'MainDataGroup',  
FILEGROUP = 'MainIndexGroup',  
FILEGROUP = 'PRIMARY'  
FROM DISK = 'W:\VulkanProd.bak'
```

```
WITH  
PARTIAL,  
REPLACE,  
RECOVERY,
```

```
MOVE N'VulkanProd' TO N'F:\Data\VulkanProd.mdf',  
MOVE N'VulkanProd_log' TO N'F:\Log\VulkanProd_log.ldf',  
MOVE N'MainDataGroup' TO N'F:\Data\MainDataGroup.ndf',  
MOVE N'MainIndexGroup' TO N'F:\Data\MainIndexGroup.ndf',  
MOVE N'FG2016' TO N'F:\Data\VulkanProdFG2016.ndf',  
MOVE N'FG2017' TO N'F:\Data\VulkanProdFG2017.ndf',  
MOVE N'FG2015' TO N'F:\Data\VulkanProdFG2015.ndf'
```

VULKAN

PRE RESTORE

Sample table: dbo.FactTablePeriode (Partition Column: PERIOD (DataType INT))

Step 7. Test Select * dbo.FactTablePeriod

One of the partitions of index 'ClusteredIndex-hiskonto' for table 'dbo.UDBHiskonto'(partition ID 72057594042843136) resides on a filegroup ("FG1994") that cannot be accessed because it is offline, restoring, or defunct. This may limit the query result.

Step 8. Test Select * dbo.FactTablePeriod WHERE PERIODE BETWEEN 201401 AND 201612

Produces rows!

Step 9. Fix Defuct Issue with **Merge**, **Switch** and **Rename**.

MERGE – Move Data from one Partition into a adjacent Partition.

UDBHis_konto



- FG2014
- (w/ FG2015 and FG2016)

```
ALTER PARTITION FUNCTION  
[PeriodFloatRange_pf]()
```

```
MERGE RANGE (201601);  
GO
```

Remember, 201601 is the lower
boundry point (RANGE)

SWITCH PARTITIONS

CREATE SHADOW TABLE

1. Create “Shadow”-table of FactTablePeriod, on FG2015. (same filegroup as FactTablePeriod, after Merge)

```
CREATE TABLE [dbo].[FactTablePeriod](  
[KON_KONTONR] [VARCHAR](11) NULL,  
[...]  
[Checksum] [INT] NULL  
) ON FG2015
```

2. Create Index on Shadow-table, aligned with Partition.

```
CREATE UNIQUE CLUSTERED INDEX [ClusteredIndexPeriode] ON [dbo].[FactTablePeriod_shadpw]  
(  
    [Period] ASC,  
    [SalesOrderID] ASC,  
    [SalesOrderDetailID] ASC  
)WITH (PAD_INDEX = OFF, STATISTICS_NORECOMPUTE = OFF, SORT_IN_TEMPDB = OFF, IGNORE_DUP_KEY = OFF,  
DROP_EXISTING = OFF, ONLINE = OFF, ALLOW_ROW_LOCKS = ON, ALLOW_PAGE_LOCKS = ON)  
GO
```

SWITCH PARTITIONS

3. Switch Partition data between Main and Shadow Tables.

```
ALTER TABLE [dbo].[FactTablePeriod]  
SWITCH PARTITION 21 TO [dbo].[FactTablePeriod_shadow]  
GO
```

4. Rename Tables

```
EXEC sp_rename 'FactTablePeriod', 'FactTablePeriod_dummy';  
  
EXEC sp_rename 'FactTablePeriod_Shadow', 'FactTablePeriod';
```

5. “Original” table name now works!

```
SELECT *  
FROM dbo.FactTablePeriod
```

VULKAN

LIVE DEMO

The Meta Data Configuration Tables

We have a total of 6 configuration tables that specifies the **core objects** and controls the **work flow** of the VULKAN system.

[vulkan].[BaseInformation] – Subsystem description

[vulkan].[BaseConfiguration] – The initial setup, per Subsystem / Server-to-server / Database

[vulkan].[DatabaseFilesConfiguration] – Information on Data- and Logfile, FileGroups, per database

[vulkan].[PartitionedTablesConfiguration] – Information concerning the partitioned tables

[vulkan].[MergeConfiguration] – Controls the workflow of the merging of Partitions on the Destination Server

[vulkan].[SwitchConfigurtation] - Controls the workflow of the Switching of Partitions on the Destination Server

[vulkan].[NonClustedIndexProcessConfiguration] – Handles the process of dropping and recreating the (Non-aligned) NonClustered Indexes on the Destination database. The system uses this table in runtime, no need to set this manually

The Meta Data Configuration Tables

[vulkan].[BaseConfiguration] – The initial setup, per Subsystem / Server-to-server / Database

[SubSystemName] - Grouping name, for a server to server backup/restore setup

[SourceServer] - Server where the Backup is created

[DestinationServer] - Server where the restore is done

[SourceDatabase] - Database to Backup

[DestinationDatabase] - Database to Restore

[RecoverModeSourceServer] - Recovery Mode of Source database (Full / Simple)

[RecoverModeDestinationServer] - Recovery Mode of Destination database (Full / Simple)

[FullOrPartialBackup] - Is it a Full (all filegroups) or Partial (selected Filegroups)

[FilePathBakFileSourceServer] - Path to folder on Source server, for .BAK file (F:\Backup\)

[FilePathBakFileDestinationServer] - Path to folder on Destination server, for .BAK file

[FilePathBakFileDestinationServerRemoteAccess] - Remote access to folder on destination server

[DeleteBackupFileOnSourceServer] - If you want to remove the BAK file on Source Server after Restore (y/n)

[DeleteBackupFileOnDestinationServer] - If you want to remove the BAK file on Destination Server after Restore (y/n)

[Active] - If the Row is Active and should be used. (y/n)

The Meta Data Configuration Tables

[vulkan].[DatabaseFilesConfiguration] – Information on Data- and Logfile, FileGroups, per database

[SubSystemName] - Grouping name, for a server to server backup/restore setup

[Database] - Name of database

[FileGroupsToBackupRestore] - One rows for each FileGroup to Backup / Restore (for Log File – “Not Applicable”)

[Type] - Datafile / LogFile

[FilePathDestination] - Specifies where on the Destination server each File should be placed

[LogicalName] - Logical name of Datafile / LogFile

[Active] - If the Row is Active and should be used (y/n)

The Meta Data Configuration Tables

[vulkan].[PartitionedTablesConfiguration] – Information concerning the partitioned tables

[SubSystemName] - Link to same SubSystem as above

[SourceServer] - Source server for the tables

[SourceDatabase] - Source database for the tables

[DestinationServer] - Destination server for the tables

[DestinationDatabase] - Destination database for the tables

[Table] - name of the table

[Active] - 'y' to use this table, 'n' to not use

The Meta Data Configuration Tables

[vulkan].[MergeConfiguration] – Controls the workflow of the merging of Partitions on the Destination Server

[SubSystemName] - Grouping name, for a server to server backup/restore setup

[DestinationServer] - Server where we are doing the Merge

[DestinationDatabase] - Database where we are doing the Merge

[PartitionFunction] - Name of function to merge

[Range] - Name of the Range to merge

[Order] - The order in which the Merges should be done

[Active] - If the Row is Active and should be used (y/n)

The Meta Data Configuration Tables

[vulkan].[SwitchConfigurtation] - Controls the workflow of the Switching of Partitions on the Destination Server

[SubSystemName] - Grouping name, for a server to server backup/restore setup

[DestinationServer] - Server where we are doing the Switch

[DestinationDatabase] - Database where we are doing the Switch

[Table] - The partitioned tables involved in the Switch

[CurrentFileGroup] - Which FileGroup are the partitioned tables Merged to? (See above)

[Active] - Active or not (y/n)

The Meta Data Configuration Tables

[vulkan].[NonClustedIndexProcessConfiguration] – Handles the process of dropping and recreating the (Non-aligned) NonClustered Indexes on the Destination database. The system uses this table in runtime, no need to set this manually

[SubSystem] - Grouping name, for a server to server backup/restore setup

[SourceServer] - Server where the Backup is created

[SourceDatabase] - Database to Backup

[DestinationServer] - Server where the restore is done

[DestinationDatabase] - Database to Restore

[CurrentSchema] – Schema for table

[CurrentTable] – Table

[RunDate] – Time of execution

[DropIndexScript] – SQL Script for Dropping Index

[CreateIndexScript] – SQL Script for Creating Index

Lets look as the tables!

The Meta Data Configuration Tables

Adding Rows to the Configuration Tables

“.....\Test_environment\VULKAN\VULKAN Dokumentasjon\VulkanAddConfigRowsProcScript.sql”

Putting it all together, BIML to the rescue!

We have six BIML scripts, that reads from Config tables and generates SSIS packages:

GeneratePreProcessPackages.biml - Generates the PreProcess Package - Saves Drop and Create sql scripts for Non-Clustered Non-Partition-Aligned Indexes for the partitioned tables in the Index meta data in [vulkan].[NonClustedIndexProcessConfiguration] table. **(This is a runtime table, no need to fill in this manually.)**

GenerateBackupPackages.biml - Generates Backup SSIS packages

GenerateRestorePackages.biml - Generates Restore SSIS packages

GenerateMergePackages.biml - Generates Merge SSIS packages

GenerateSwitchPackages.biml - Generates Switch SSIS packages

GeneratePostProcessPackages.biml - Recreates the Non-Clustered Non-Partition-Aligned Indexes for the partitioned tables, on the Destination server.To generate the SSIS packages, aligned with the Config tables, right click each Biml-file and select: Generate SSIS packages. **Remember, if you add/edit rows in ConFig Tables, you have re-generate the affected SSIS package, and Deploy this to the SSIS server.**

VULKAN

LIVE DEMO